

YData: Introduction to Data Science



Class 05: Array computation continued

Overview

Brief review of:

- Statistics and visualizations of quantitative data

Continuation of numpy:

- Review: NumPy arrays
- Numerical array computations
- Higher dimensional numerical arrays
- Image manipulation

If there is time:

- Introduction to pandas Series and DataFrames



Announcement: Homework 2

Homework 2 has been posted!

It is due on Gradescope on **Sunday February 4th at 11pm**

- **Be sure to mark each question on Gradescope!**

Notes:

- On problem 3, if the images are not showing up make sure to run the cells where the images are embedded
 - If you figure it out, help other people on Ed!
- I think the problems are at a reasonable level of difficulty
 - I will get your feedback...

How would describe the pace of the class so far?

Way too slow		0 %	✓
Too slow	3 respondents	3 %	
About right	71 respondents	75 %	
Too fast	20 respondents	21 %	
Way too fast	1 respondent	1 %	



Questions?

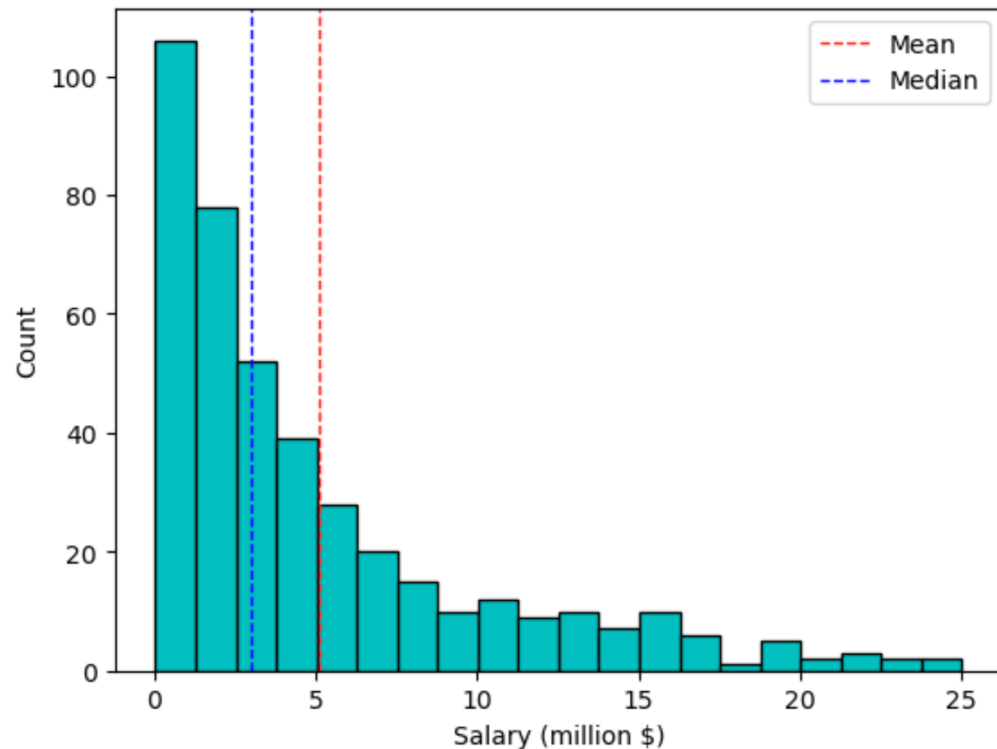


Review of array computations

Very quick review: quantitative data

We can visualize quantitative data using histograms

Two statistics that measure “central value” are the *mean* and the *median*



```
import matplotlib.pyplot as plt
```

```
plt.hist(data)
```

```
plt.xlabel("salary (million $)")
```

```
plt.ylabel("count")
```

```
import statistics
```

```
statistics.mean(data)
```

```
statistics.median(data)
```

Very quick review: outliers

Q: What is an **outlier**?

- A: An observed value that is notably distinct from the other values in a dataset

Q: Why are they problematic?

- A: can potentially have a large influence on the statistics you calculate

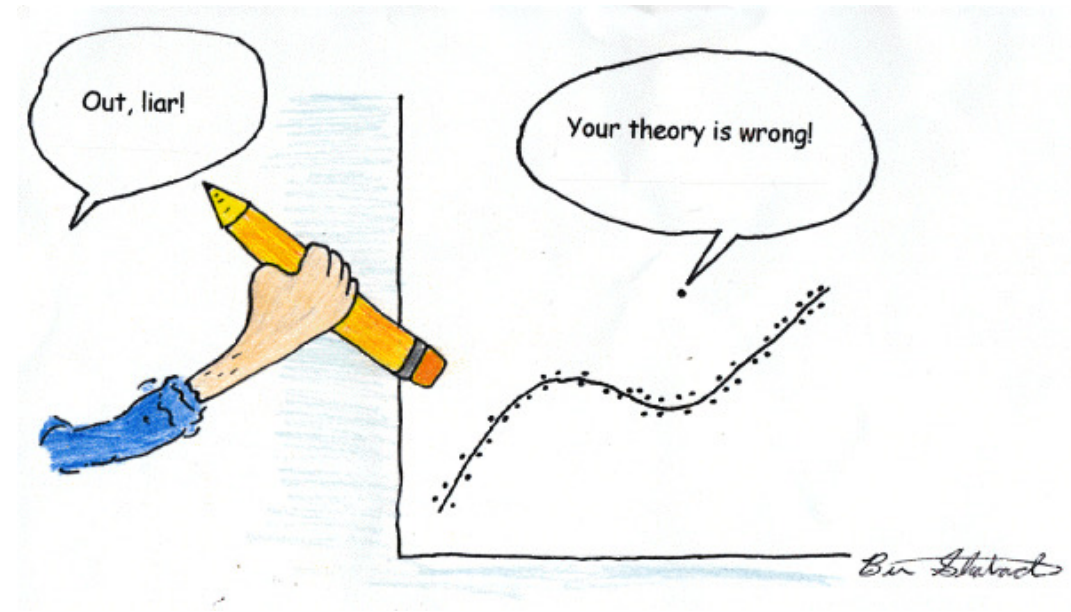
Q: What should you do if you have an outlier in your data?

A: See if you can understand what is causing it!

- If it's an error, delete the point
- If it's a real value, make sure it is not having a big effect on your conclusions, and/or use resistant statistics

Q: Is the mean and/or median resistant?

- A: The median is resistant while the mean is not



Review of NumPy arrays and functions

Processing data that is all of the same type can be more efficient than processing data of mixed types

The *NumPy package* stores and processes data that is all of the same type using *ndarray*

The package also contains functions that operate efficiently on these arrays



Review: ndarrays

Let's review this in Jupyter!

```
import numpy as np
```

```
my_array = np.array([1, 2, 3]) # creating an ndarray from a list
```

```
my_array[0] # accessing the 0th element of the ndarray
```

```
my_array.dtype # get the type of elements stored in the array
```

```
my_array.shape # get the dimension of the array
```

```
my_array.astype('str') # convert the numbers to strings
```

```
sequential_nums = np.arange(1, 10) # creates numbers 1 to 9
```

NumPy functions on numerical arrays

The NumPy functions:

- `np.sum()`
- `np.max()`, `np.min()`
- `np.mean()`, `np.median()`
- `np.diff()` # takes the difference between elements
- `np.cumsum()` # cumulative sum

There are also "broadcast" functions that operate on all elements in an array

- `my_array = np.array([12, 4, 6, 3, 4, 3, 7, 4])`
- `my_array * 2`

- `my_array2 = np.array([10, 9, 2, 8, 9, 3, 8, 5])`
- `my_array - my_array2`

Let's explore this in Jupyter!

Boolean arrays

It is often to compare all values in an ndarray to a particular value

- `my_array = np.array([12, 4, 6, 3, 4, 3, 7, 4])`
- `my_array < 5` # any guesses what this will return
 - `array([False, True, False, True, True, True, False, True])`

This can be useful for calculating proportions

- `True == 1` and `False == 0`
- Taking the sum of a Boolean array gives the total number of `True` values
- The number of `True` 's divided by the length is the proportion
 - Or we can use the `np.mean()` function

Categorical Variable

PLAYER	POSITION	TEAM	SALARY
str	str	str	f64
"Paul Millsap"	"PF"	"Atlanta Hawks"	18.671659
"Al Horford"	"C"	"Atlanta Hawks"	12.0
"Tiago Splitter..."	"C"	"Atlanta Hawks"	9.75625
"Jeff Teague"	"PG"	"Atlanta Hawks"	8.0
"Kyle Korver"	"SG"	"Atlanta Hawks"	5.746479

$$\text{Proportion centers} = \frac{\text{number of centers}}{\text{total number}}$$

Let's explore this in Jupyter!

Boolean masking

We can also use Boolean arrays to return values in another array

- This is called "Boolean masking" or "Boolean indexing"

```
my_array = np.array([12, 4, 6, 3, 1])  
boolean_mask = np.array([False, True, False, True, True])  
  
smaller_array = my_array[boolean_mask]
```

This can be useful for calculating statistics on data that meet particular criteria:

- `np.mean(my_array[my_array < 5])` # what does this do?

Boolean masking

Suppose you wanted to get the average salary of NBA players who were centers

If you had these two ndarrays:

- **Position:** The position of all NBA players
- **Salary:** Their salaries

Could you do it?



Let's explore this in Jupyter!

Higher dimensional arrays

We can make higher dimensional arrays

- (matrices and tensors)

```
my_matrix = np.array([1, 2, 3], [4, 5, 6], [7, 8, 9])
```

```
my_matrix
```

We can slice higher dimensional array

- `my_matrix[0:2, 0:2]`

We can apply operations to rows, columns, etc.

- `np.sum(my_matrix, axis = 0)` # sum the values down rows

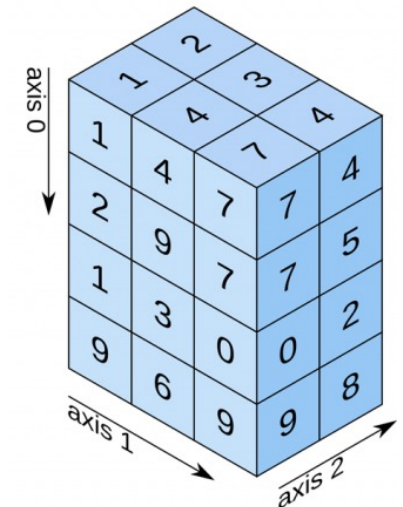
Let's explore this in Jupyter!

2D array



shape: (2, 3)

3D array



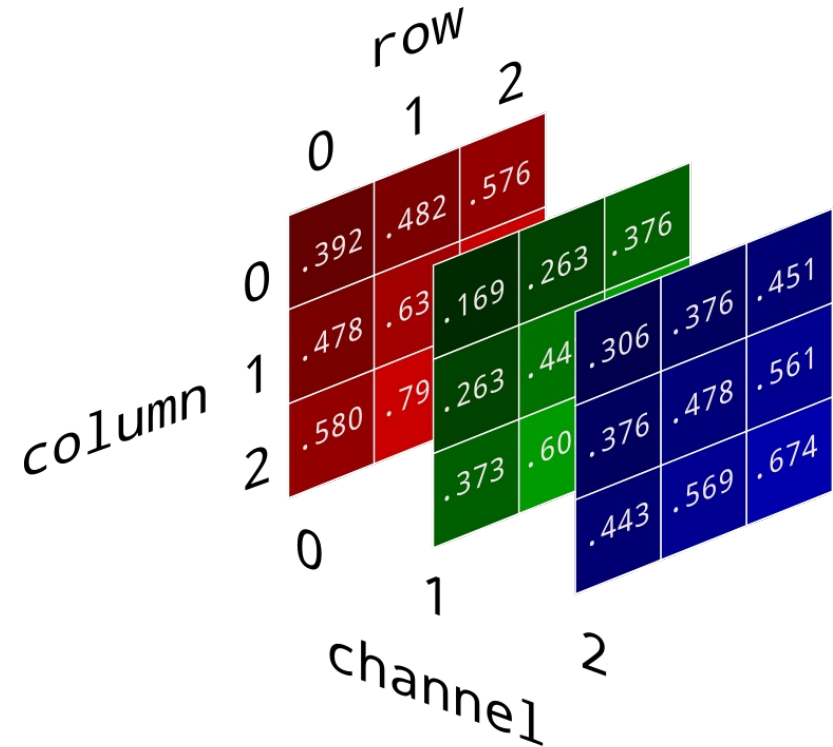
shape: (4, 3, 2)

Image processing

3-dimensional numerical arrays are often used to store digital images

- RGB image = Red, Green, Blue matrices

We can use masking and other array operations to process images



Let's explore this in Jupyter!



Series and Tables

Pandas: Series and DataFrames

“[pandas](#) is an open source, BSD-licensed library providing high-performance, easy-to-use data structures and data analysis tools for the Python programming language.”



There are two main data structures in pandas:

- **Series:** represent one-dimensional data
- **DataFrames:** represent data tables
 - i.e., relational data

pandas Series

pandas Series are: One-dimensional ndarray with axis labels

- (including time series)

Example: egg_prices

DATE

1980-01-01 0.879

1980-02-01 0.774

1980-03-01 0.812

Index



values



pandas Series

We can access elements by Index ***name*** using **.loc**

- `egg_prices.loc["1980-01-01"]`

We can access elements by Index ***number*** using **.iloc**

- `egg_prices.iloc[0]`

Let's explore this in Jupyter!

pandas DataFrames

Pandas DataFrame hold
Table data

This is one of the most
useful formats to extract
insights from datasets

Often we read data into
a DataFrame using:

- `pd.read_csv("file.csv")`

Variables

Cases

PLAYER	POSITION	TEAM	SALARY
str	str	str	f64
"Paul Millsap"	"PF"	"Atlanta Hawks"	18.671659
"Al Horford"	"C"	"Atlanta Hawks"	12.0
"Tiago Splitter..."	"C"	"Atlanta Hawks"	9.75625
"Jeff Teague"	"PG"	"Atlanta Hawks"	8.0
"Kyle Korver"	"SG"	"Atlanta Hawks"	5.746479

Let's explore this in Jupyter!

Next class: pandas continued...

