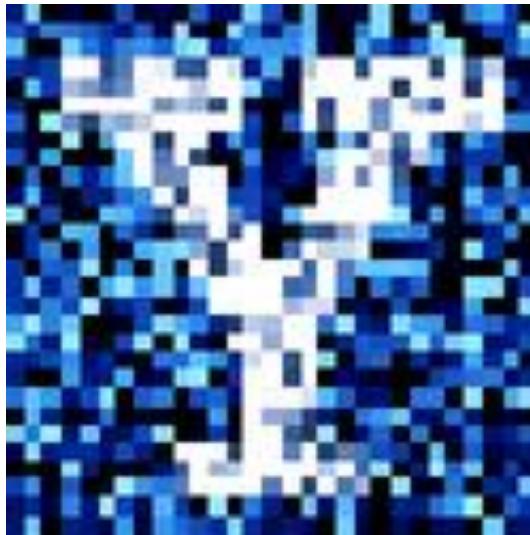


# YData: Introduction to Data Science



Lecture 24: Linear regression and unsupervised learning

# Overview

Quick review of classification and the KNN classifier

Feature normalization

Linear regression

Unsupervised learning/clustering

- K-means cluster
- If there is time: hierarchical clustering

# Project timeline

~~Sunday, November 16<sup>th</sup>~~

- Projects are due on Gradescope at 11pm
- Email a pdf of your project to your peer reviewers
  - A list of whose paper you will review is posted on Canvas
  - Fill out the draft reflection on Canvas



Sunday, November 23<sup>rd</sup>

- A template for doing your review has been posted
- Jupyter notebook files with your reviews need to be emailed to the authors
- A pdf containing all three reviews needs to be uploaded Gradescope

Sunday, December 7<sup>th</sup>

- Project is due on Gradescope
- Add the peer reviews of your project to the Appendix of your project

Homework 9 has been posted

- It is due **Monday** December 1<sup>st</sup>

# Prediction: regression and classification

We “learn” a function  $f$

- $f(\mathbf{x}) \longrightarrow y$

Input:  $\mathbf{x}$  is a data vector of "features"

Output:

- Regression: output is a real number ( $y \in \mathbb{R}$ )
- Classification: output is a categorical variable  $y_k$


# Training a classifier



# K-Nearest Neighbor Classifier (KNN)

**Training the classifier:** Store all the features with their labels

**Making predictions:** The label of closest k training points is returned




# KNN classifiers using scikit-learn

We can fit and evaluate the performance of a KNN classifier using:

```
knn = KNeighborsClassifier(n_neighbors = 1)      # construct a classifier  
  
knn.fit(X_features, y_labels)      # train the classifier  
  
penguin_predictions = knn.predict(X_penguin_features) # make predictions  
  
np.mean(penguin_predictions == y_penguin_labels)      # get accuracy
```

# Review: overfitting

Overfitting occurs when our classifier matches too close to the training data and doesn't capture the true underlying patterns




[Overfitting song...](#)

To avoid overfitting we split our data into a training and test set

- Classifier learns relationship on training data
- Classifier's performance is evaluated on the test data



We can use k-fold cross-validation to get a better estimate of the test accuracy



# Review: Other classifiers

We also built our own KNN classifier



Scikit-learn makes it easy to try out different classifiers get their cross-validation performance

- E.g., SVM, random forests, neural networks, etc.

`svm = LinearSVC()`

```
scores = cross_val_score(svm,  
X_features, y_labels, cv = 5)
```

```
scores.mean()
```



# Review: steps to build a KNN classifier

We build our KNN classifier by creating a series of functions...

## 1. `euclid_dist(x1, x2)`

$$D = \sqrt{(x_0 - x_1)^2 + (y_0 - y_1)^2 + (z_0 - z_1)^2}$$

- Calculates the Euclidean distance between two points  $x_1$  and  $x_2$

## 2. `get_labels_and_distances(test_point, X_train_features, y_train_labels)`

- Finds the distance between a test point and all the training points

## 3. `classify_point(test_point, k, X_train_features, y_train_labels)`

- Classifies one test point by returning the majority label of the  $k$  closest points

## 4. `classify_all_test_data(X_test_data, k, X_train_features, y_train_labels)`

- Classifiers all test points

# Feature normalization

# Review: Distance between two points

Two features x and y:  $D = \sqrt{(x_0 - x_1)^2 + (y_0 - y_1)^2}$

Three features x, y, and z:  $D = \sqrt{(x_0 - x_1)^2 + (y_0 - y_1)^2 + (z_0 - z_1)^2}$

- And so on for more features...

It's important the features are standardized

- If not, features that typically have larger values will dominate the distance measurement

# Feature normalization

In order to deal with features that are measured on very different scales, we can normalize the features

With a z-score normalization, we normalize each feature to have:

- A mean of 0
- A standard deviation of 1

We can do this in Python using:

```
scalar = StandardScaler()  
scalar.fit(X_train)  
X_train_transformed = scalar.transform(X_train)  
X_test_transformed = scalar.transform(X_test)
```

bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g
46.1	15.1	215.0	5100.0
37.3	17.8	191.0	3350.0
51.3	18.2	197.0	3750.0
39.5	16.7	178.0	3250.0
48.7	15.1	222.0	5350.0

$$\bar{x}$$

$$s$$

$$z_i = \frac{x_i - \bar{x}}{s}$$



# Feature normalization

To avoid overfitting ("data leakage") we can:

- Calculate the mean and standard deviation on the training
- Apply these means and standard deviations to normalize the training and test sets

To do this in a cross-validation loop, we can use a pipeline:

```
scalar = StandardScaler()
```

```
knn = KNeighborsClassifier(n_neighbors = 1)
```

```
cv = KFold(n_splits=5)
```

```
pipeline = Pipeline([('transformer', scalar), ('estimator', knn)])
```

```
scores = cross_val_score(pipeline, X_penguin_features, y_penguin_labels, cv = cv)
```

Let's explore this in Jupyter!

# Linear regression

# Prediction: regression and classification

We “learn” a function  $f$

- $f(x) \rightarrow y$

Input:  $x$  is a data vector of "features"




Output:

- Regression: output is a real number ( $y \in \mathbb{R}$ )
- Classification: output is a categorical variable  $y_k$

# Regression

Regression is method of using one variable  $x$  to predict the value of a second variable  $y$

- i.e.,  $\hat{y} = f(x)$
- Linear regression:  $\hat{y} = \text{intercept} + \text{slope} \cdot x$   
$$\hat{y} = b_0 + b_1 \cdot x$$



The coefficients for these regression models are found by minimizing root mean square error (RMSE)

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

[Regression line app](#)

# Multiple regression


In multiple regression we try to predict a quantitative response variable  $y$  using several features  $x_1, x_2, \dots, x_k$

We estimate coefficients using a data set to make predictions  $\hat{y}$

$$\hat{y} = b_0 + b_1 \cdot x_1 + b_2 \cdot x_2 + \dots + b_k \cdot x_k$$



Learn the  $b_i$ 's on the training set. Assess prediction accuracy on test set.




# Multiple regression

$$\hat{y} = b_0 + b_1 \cdot x_1 + b_2 \cdot x_2 + \dots + b_k \cdot x_k$$

There are many uses for multiple regression models including:

- To make predictions as accurately as possible
- To understand which predictors ( $x$ ) are related to the response variable ( $y$ )



# Linear regression models in scikit-learn

We can use scikit-learn to create linear regression models

- You can also use the stats models package to do this

```
linear_model = LinearRegression() # construct a linear regression model
```

```
linear_model.fit(X_train_features, y_train) # train the classifier
```

"Learns"  $b_0, b_1, \dots$  by minimizing  
the RMSE on the training data

# train the classifier


Numeric values  
to predict

```
y_predictions = linear_model.predict(X_test_features) # make predictions
```

```
RMSE = np.sqrt(np.mean((y_test - y_predictions)**2)) # get the RMSE
```

# Real world example

Rather than predicting stock prices...




Let's try this in Jupyter!

# Inference for simple linear regression

# Inference for simple linear regression



The Greek letter  $\beta$  is used to denote the slope *of the population*

The letter  $b$  is used to denote the slope *of the sample*



Population:  $\beta_0$  and  $\beta_1$



Sample estimates:  $b_0$  and  $b_1$



# Hypothesis test for regression coefficients

We can run a hypothesis test to see if there is a linear relationship between a feature  $x$ , and a response variable  $y$

There is only one [hypothesis test!](#)



# Hypothesis test for regression coefficients


We can run a hypothesis test to see if there is a linear relationship between a feature  $x$ , and a response variable  $y$

- $H_0: \beta_1 = 0$  (slope is 0, so no linear relationship between  $x$  and  $y$ )
- $H_A: \beta_1 \neq 0$

One type of hypothesis test we can run is based on a t-statistic:

- The t-statistic comes from a t-distribution with  $n - k$  degrees of freedom
  - (If a few conditions are met)

We could also shuffle our data, fit models, and extract  $b_1$  coefficients to create a null distribution



# Confidence intervals for regression coefficients

The confidence interval for the slope coefficients:

$$\beta_1$$


We can use the statsmodels package to run hypothesis tests and create confidence intervals using “parametric” methods based on t-distributions

```
import statsmodels.api as sm

sm_linear_model = sm.OLS(y_train,
                          X_train_with_constant).fit()

sm_linear_model.summary()
```

Let's try this in Jupyter!



# Unsupervised learning

# Supervised learning and unsupervised learning

In **supervised learning** we have a set of features X, along with a label y


- We use the features X to predict y on new data

In **unsupervised learning**, we have features X, but **no** response variable y


Unsupervised learning can be useful in order to find structure in the data and to visualize patterns

A key challenge in unsupervised learning is that there is no real ground truth response variable y

- So we don't have measures like the mean prediction accuracy



**Supervised Learning**



**Unsupervised Learning**


# Unsupervised learning

Given we are almost at the end of the semester, we will focus on clustering, which is one type of unsupervised learning


In **clustering** we try to group similar data points together

Another type of unsupervised learning:

- **Dimensionality reduction** where we try to find a smaller set of features that captures most of the variability original larger feature set
  - E.g., principal component analysis (PCA)



**Supervised Learning**



**Unsupervised Learning**

# Clustering

# Clustering

Clustering divides  $n$  data points  $x_i$ 's into subgroups

- Data points in the same group are similar/homogeneous
- Data points in different groups are different from each other

A diagram illustrating a data matrix. A red bracket on the left indicates the number of rows is  $n$ . A red bracket at the top indicates the number of columns is  $p$ . The matrix itself consists of four horizontal rows. The first row is highlighted with a blue border and contains entries  $x_{11}, x_{12}, \dots, x_{1p}$ . The second row is also highlighted with a blue border and contains entries  $x_{21}, x_{22}, \dots, x_{2p}$ . The third row contains three vertical ellipses between the second and third columns, indicating intermediate data points. The fourth row is highlighted with a red border and contains entries  $x_{n1}, x_{n2}, \dots, x_{np}$ .

$x_{11}$	$x_{12}$	$\cdots$	$x_{1p}$
$x_{21}$	$x_{22}$	$\cdots$	$x_{2p}$
$\vdots$	$\vdots$	$\ddots$	$\vdots$
$x_{n1}$	$x_{n2}$	$\cdots$	$x_{np}$

Examples:

- Examining gene expression levels to group cancer types together
- Examining consumer purchasing behavior to perform market segmentation

Clustering can be:

- **Flat:** no structure beyond dividing points into groups
- **Hierarchical:** Population is divided into smaller and smaller groups (tree like structure)

# K-means clustering

K-means clustering partitions the data into  $K$  distinct, non-overlapping clusters

- i.e., each data point  $x_i$  belongs to exactly one cluster  $C_k$

The number of clusters,  $K$ , needs to be specified prior to running the algorithm


The goal is to minimize the within-cluster variation

- e.g., to make the Euclidean distance for all points within a cluster as small as possible

Finding the exact optimal solution is computationally intractable (there are  $k^n$  possible partitions), but a simple algorithm exists to find a local optimum which is often works well in practice

# K-means clustering

1. Randomly assign points to clusters  $C_k$
2. Calculate cluster centers as means of points in each cluster
3. Assign points to the closest cluster center
4. Recalculate cluster center as the mean of points in each cluster
5. Repeat steps 3 and 4 until convergence



# K-means clustering

Because only a local minimum is found, different random initializations will lead to different solutions

- One should run the algorithm multiple times to get better solutions

Let's explore this in Jupyter!




# Hierarchical clustering

# Hierarchical clustering

In **hierarchical clustering** we create a dendrogram which is a tree-based representation of successively larger clusters.

We can cut the dendrogram at any point to create as many clusters as desired


- i.e., don't need to specify the number of clusters,  $K$ , beforehand



# Hierarchical clustering



We can create a hierarchical clustering of the data using simple bottom-up agglomerative algorithm:

1. Choosing a (dis)similarity measure
  - E.g., The Euclidean distance
2. Initializing the clustering by treating each point as its own cluster
3. Successively merging the pair of clusters that are most similar
  - i.e., calculate the similarity between all pairs of clusters and merging the pair that is most similar
4. Stopping when all points have been merged into a single cluster



# Hierarchical clustering

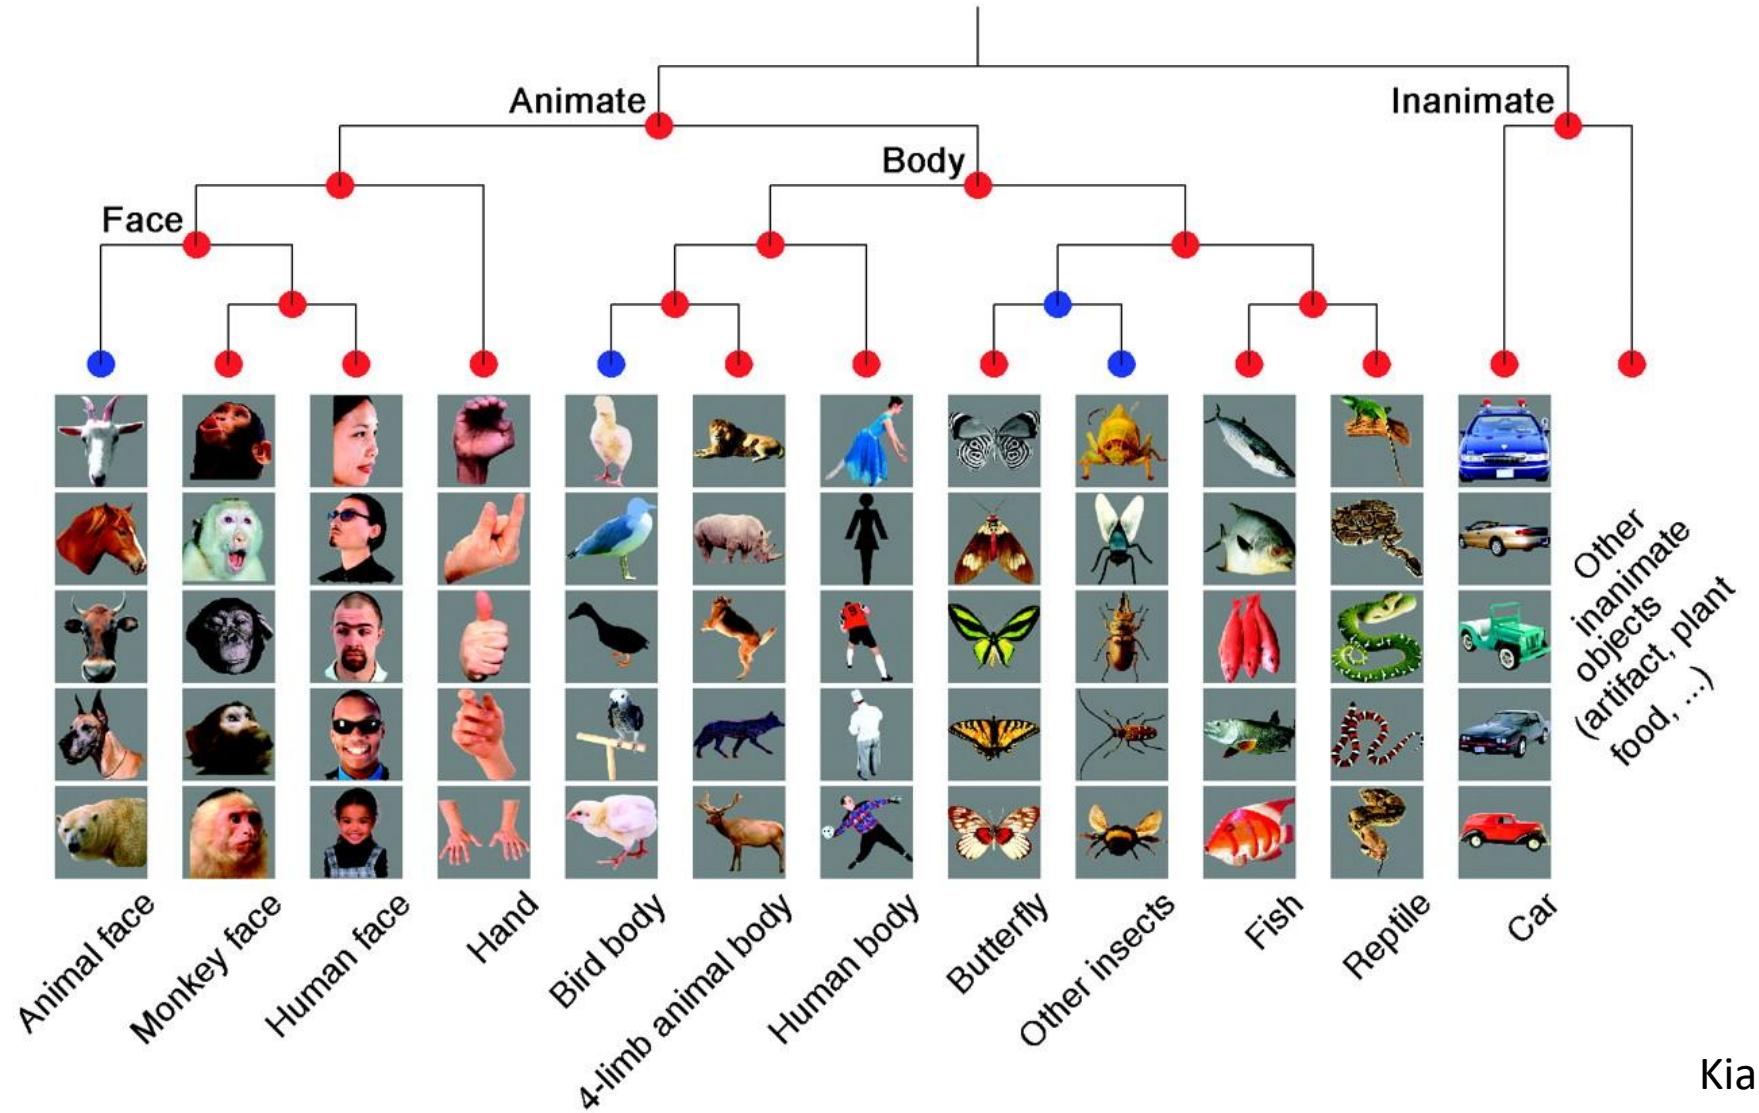
The vertical height that two clusters/points merge show how similar the two *clusters* are



Note: horizontal distance between *individual points* is not important:

- point 9 is considered as similar to point 2 as it is to point 7


# Hierarchical clustering example



# Issues with clustering

Choices made can effect the results:

- Feature normalization and/or dissimilarity measure
- K-means: choice of K
- For hierarchical cluster: linkage and cut height



**Unsupervised Learning**

Potential approaches to deal with these issues:

- Try a few methods and see if one gives interesting/useful results
- Validate that you get similar results on a second set of data

Let's explore this in Jupyter!